

Bayesian networks II

Advanced Artificial Intelligence: Lecture

Heriberto Cuayáhuitl

<http://staff.lincoln.ac.uk/hcuayahuitl>

School of Computer Science, University of Lincoln



UNIVERSITY OF
LINCOLN

16 October 2019

Overview

- ▶ Last Week: Exact Inference in Bayes Nets
 - ▶ Inference by Enumeration
 - ▶ Inference by Variable Elimination
- ▶ This Week: Approximate Inference in Bayes Nets
 - ▶ Rejection Sampling
 - ▶ Likelihood Weighting
 - ▶ Gibbs Sampling

Reading of last week: Russell & Norvig, (2013). Artificial Intelligence: A Modern Approach, **chapter 14** (until 14.4)

Reading of this week: Russell & Norvig, (2013). Artificial Intelligence: A Modern Approach, **chapter 14** (from 14.5)

Probabilistic Sampling


- ▶ **Idea:** Draw N samples from prior/conditional distributions and compute approximate posterior probabilities
- ▶ **Why?** Approximate inference methods give reasonable answers and can be applied to large probability models
- ▶ **Example:** Sampling 8 times from $P(C=\text{red}) = 0.6$, $P(C=\text{green}) = 0.1$, and $P(C=\text{blue}) = 0.3$,



gives the following distribution: $P(C) = \langle 5/8, 1/8, 2/8 \rangle$

Sampling in Bayesian Networks

Sampling=generation of samples from a known probability distrib.



Assume topological order [*Cloudy*, *Sprinkler*, *Rain*, *WetGrass*]:

- ▶ Sample from $\mathbf{P}(\textit{Cloudy}) = \langle 0.5, 0.5 \rangle$
- ▶ Sample from $\mathbf{P}(\textit{Sprinkler} | \textit{Cloudy} = \textit{true}) = \langle 0.1, 0.9 \rangle$
- ▶ Sample from $\mathbf{P}(\textit{Rain} | \textit{Cloudy} = \textit{true}) = \langle 0.8, 0.2 \rangle$
- ▶ Sample .. $\mathbf{P}(\textit{WetGrass} | \textit{Sprinkler} = \textit{false}, \textit{Rain} = \textit{true}) = \langle 0.9, 0.1 \rangle$

Prior Sampling

- ▶ This function returns events like $[true, false, true, true]$ — assuming the Bayesian network and variable ordering above

function PRIOR-SAMPLE(bn) **returns** an event sampled from the prior specified by bn

inputs: bn , a Bayesian network specifying joint distribution $\mathbf{P}(X_1, \dots, X_n)$

$\mathbf{x} \leftarrow$ an event with n elements

foreach variable X_i **in** X_1, \dots, X_n **do**

$\mathbf{x}[i] \leftarrow$ a random sample from $\mathbf{P}(X_i \mid \text{parents}(X_i))$

return \mathbf{x}

Prior Sampling

- ▶ This function returns events like $[true, false, true, true]$ — assuming the Bayesian network and variable ordering above

function PRIOR-SAMPLE(bn) **returns** an event sampled from the prior specified by bn

inputs: bn , a Bayesian network specifying joint distribution $\mathbf{P}(X_1, \dots, X_n)$

$\mathbf{x} \leftarrow$ an event with n elements

foreach variable X_i **in** X_1, \dots, X_n **do**

$\mathbf{x}[i] \leftarrow$ a random sample from $\mathbf{P}(X_i \mid \text{parents}(X_i))$

return \mathbf{x}

- ▶ The probability of an event can be estimated as

$$P(x_1, \dots, x_n) \approx \frac{N_{PS}(x_1, \dots, x_n)}{N}$$

Prior Sampling

- ▶ This function returns events like $[true, false, true, true]$ — assuming the Bayesian network and variable ordering above

function PRIOR-SAMPLE(bn) **returns** an event sampled from the prior specified by bn

inputs: bn , a Bayesian network specifying joint distribution $\mathbf{P}(X_1, \dots, X_n)$

$\mathbf{x} \leftarrow$ an event with n elements

foreach variable X_i **in** X_1, \dots, X_n **do**

$\mathbf{x}[i] \leftarrow$ a random sample from $\mathbf{P}(X_i \mid \text{parents}(X_i))$

return \mathbf{x}

- ▶ The probability of an event can be estimated as
$$P(x_1, \dots, x_n) \approx \frac{N_{PS}(x_1, \dots, x_n)}{N}$$
- ▶ **Example** : generating 1000 samples from the sprinkler network, and observing that 511 of them have $Rain = true$, we get the estimated probability $\hat{P}(Rain = true) = 0.511$

Rejection Sampling: Concepts and Example

- ▶ Can be used to compute conditional probabilities of the form $\mathbf{P}(X|\mathbf{e})$, given query variable X and evidence \mathbf{e}

Rejection Sampling: Concepts and Example

- ▶ Can be used to compute conditional probabilities of the form $\mathbf{P}(X|\mathbf{e})$, given query variable X and evidence \mathbf{e}
- ▶ **Idea**: to generate samples from the prior distribution specified by the net and reject all those that do not match the evidence

Rejection Sampling: Concepts and Example

- ▶ Can be used to compute conditional probabilities of the form $\mathbf{P}(X|\mathbf{e})$, given query variable X and evidence \mathbf{e}
- ▶ **Idea**: to generate samples from the prior distribution specified by the net and reject all those that do not match the evidence
- ▶ Let $\hat{\mathbf{P}}(X|\mathbf{e}) = \frac{\mathbf{N}_{PS}(X, \mathbf{e})}{N_{PS}(\mathbf{e})} = \alpha \mathbf{N}_{PS}(X, \mathbf{e}) \approx \frac{\mathbf{P}(X, \mathbf{e})}{P(\mathbf{e})} = \mathbf{P}(X|\mathbf{e})$

Rejection Sampling: Concepts and Example

- ▶ Can be used to compute conditional probabilities of the form $\mathbf{P}(X|\mathbf{e})$, given query variable X and evidence \mathbf{e}
- ▶ **Idea**: to generate samples from the prior distribution specified by the net and reject all those that do not match the evidence
- ▶ Let $\hat{\mathbf{P}}(X|\mathbf{e}) = \frac{\mathbf{N}_{PS}(X, \mathbf{e})}{N_{PS}(\mathbf{e})} = \alpha \mathbf{N}_{PS}(X, \mathbf{e}) \approx \frac{\mathbf{P}(X, \mathbf{e})}{P(\mathbf{e})} = \mathbf{P}(X|\mathbf{e})$
- ▶ **Example**: estimate $\hat{\mathbf{P}}(\text{Rain}|\text{Sprinkler}=\text{true})$ from 100 samples
 - ▶ 73 samples have *Sprinkler=false* (rejected)
 - ▶ 27 samples: 8 have *Rain=true* and 19 have *Rain=false*
 - ▶ $\hat{\mathbf{P}}(\text{Rain}|\text{Sprinkler}=\text{true}) \approx \alpha \langle 8, 19 \rangle = \langle 0.296, 0.704 \rangle$

Rejection Sampling: Algorithm

function REJECTION-SAMPLING(X, \mathbf{e}, bn, N) **returns** an estimate of $\mathbf{P}(X|\mathbf{e})$

inputs: X , the query variable

\mathbf{e} , observed values for variables \mathbf{E}

bn , a Bayesian network

N , the total number of samples to be generated

local variables: \mathbf{N} , a vector of counts for each value of X , initially zero

for $j = 1$ to N **do**

$\mathbf{x} \leftarrow \text{PRIOR-SAMPLE}(bn)$

if \mathbf{x} is consistent with \mathbf{e} **then**

$\mathbf{N}[x] \leftarrow \mathbf{N}[x] + 1$ where x is the value of X in \mathbf{x}

return NORMALIZE(\mathbf{N})

Likelihood Weighting: Concepts

- ▶ **Motivation:** Rejection sampling rejects so many samples!

Likelihood Weighting: Concepts

- ▶ **Motivation:** Rejection sampling rejects so many samples!
- ▶ **Key ideas:**
 1. It avoids rejecting samples by generating only events that are consistent with evidence \mathbf{e} .
 2. It fixes the values for the evidence variables \mathbf{E} .
 3. It samples only the non-evidence variables.

Likelihood Weighting: Example¹

Query: $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}, \text{WetGrass} = \text{true})$

1. $w \leftarrow 1$

¹Assume topological order [*Cloudy*, *Sprinkler*, *Rain*, *WetGrass*]

Likelihood Weighting: Example¹

Query: $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}, \text{WetGrass} = \text{true})$

1. $w \leftarrow 1$
2. $w \leftarrow w \times P(\text{Cloudy} = \text{true}) = 1 \times 0.5 = 0.5$

¹Assume topological order [*Cloudy*, *Sprinkler*, *Rain*, *WetGrass*]

Likelihood Weighting: Example¹

Query: $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}, \text{WetGrass} = \text{true})$

1. $w \leftarrow 1$
2. $w \leftarrow w \times P(\text{Cloudy} = \text{true}) = 1 \times 0.5 = 0.5$
3. Sample from $\mathbf{P}(\text{Sprinkler} | \text{Cloudy} = \text{true}) = \langle 0.1, 0.9 \rangle$ because $\text{Sprinkler} \notin \mathbf{E}$; assume this returns *false*

¹Assume topological order [*Cloudy*, *Sprinkler*, *Rain*, *WetGrass*]

Likelihood Weighting: Example¹

Query: $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}, \text{WetGrass} = \text{true})$

1. $w \leftarrow 1$
2. $w \leftarrow w \times P(\text{Cloudy} = \text{true}) = 1 \times 0.5 = 0.5$
3. Sample from $\mathbf{P}(\text{Sprinkler} | \text{Cloudy} = \text{true}) = \langle 0.1, 0.9 \rangle$ because $\text{Sprinkler} \notin \mathbf{E}$; assume this returns *false*
4. Sample from $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}) = \langle 0.8, 0.2 \rangle$ because $\text{Rain} \notin \mathbf{E}$; assume this returns *true*

¹Assume topological order [*Cloudy*, *Sprinkler*, *Rain*, *WetGrass*]

Likelihood Weighting: Example¹

Query: $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}, \text{WetGrass} = \text{true})$

1. $w \leftarrow 1$
2. $w \leftarrow w \times P(\text{Cloudy} = \text{true}) = 1 \times 0.5 = 0.5$
3. Sample from $\mathbf{P}(\text{Sprinkler} | \text{Cloudy} = \text{true}) = \langle 0.1, 0.9 \rangle$ because $\text{Sprinkler} \notin \mathbf{E}$; assume this returns *false*
4. Sample from $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}) = \langle 0.8, 0.2 \rangle$ because $\text{Rain} \notin \mathbf{E}$; assume this returns *true*
5. $w \leftarrow w \times P(\text{WetGrass} = \text{true} | \text{Sprinkler} = \text{false}, \text{Rain} = \text{true}) = 0.5 \times 0.9 = 0.45$

¹Assume topological order [*Cloudy*, *Sprinkler*, *Rain*, *WetGrass*]

Likelihood Weighting: Example¹

Query: $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}, \text{WetGrass} = \text{true})$

1. $w \leftarrow 1$
2. $w \leftarrow w \times P(\text{Cloudy} = \text{true}) = 1 \times 0.5 = 0.5$
3. Sample from $\mathbf{P}(\text{Sprinkler} | \text{Cloudy} = \text{true}) = \langle 0.1, 0.9 \rangle$ because $\text{Sprinkler} \notin \mathbf{E}$; assume this returns *false*
4. Sample from $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{true}) = \langle 0.8, 0.2 \rangle$ because $\text{Rain} \notin \mathbf{E}$; assume this returns *true*
5. $w \leftarrow w \times P(\text{WetGrass} = \text{true} | \text{Sprinkler} = \text{false}, \text{Rain} = \text{true}) = 0.5 \times 0.9 = 0.45$
6. Return event $[\text{true}, \text{false}, \text{true}, \text{true}]$ with weight 0.45

¹Assume topological order $[\text{Cloudy}, \text{Sprinkler}, \text{Rain}, \text{WetGrass}]$

Likelihood Weighting: Algorithm

function LIKELIHOOD-WEIGHTING(X, \mathbf{e}, bn, N) **returns** an estimate of $\mathbf{P}(X|\mathbf{e})$
 inputs: X , the query variable
 \mathbf{e} , observed values for variables \mathbf{E}
 bn , a Bayesian network specifying joint distribution $\mathbf{P}(X_1, \dots, X_n)$
 N , the total number of samples to be generated
 local variables: \mathbf{W} , a vector of weighted counts for each value of X , initially zero

 for $j = 1$ to N **do**
 $\mathbf{x}, w \leftarrow \text{WEIGHTED-SAMPLE}(bn, \mathbf{e})$
 $\mathbf{W}[x] \leftarrow \mathbf{W}[x] + w$ where x is the value of X in \mathbf{x}
 return NORMALIZE(\mathbf{W})

function WEIGHTED-SAMPLE(bn, \mathbf{e}) **returns** an event and a weight

 $w \leftarrow 1$; $\mathbf{x} \leftarrow$ an event with n elements initialized from \mathbf{e}
 foreach variable X_i **in** X_1, \dots, X_n **do**
 if X_i is an evidence variable with value x_i in \mathbf{e}
 then $w \leftarrow w \times P(X_i = x_i \mid \text{parents}(X_i))$
 else $\mathbf{x}[i] \leftarrow$ a random sample from $\mathbf{P}(X_i \mid \text{parents}(X_i))$
 return \mathbf{x}, w

Gibbs Sampling: Concepts

- Motivation:

While Likelihood Weighting can be more efficient than Rejection sampling, its performance degrades as the number of evidence variables increases (e.g. due to low weights)

Gibbs Sampling: Concepts

- Motivation:


While Likelihood Weighting can be more efficient than Rejection sampling, its performance degrades as the number of evidence variables increases (e.g. due to low weights)

- Key ideas:

1. Instead of generating each sample from scratch, MCMC (Markov Chain Monte Carlo) algorithms generate each sample by making a random change to the preceeding sample.
2. Start with an arbitrary event (called *state*), with evidence \mathbf{e} fixed, and generate a next state by sampling a value from one of the non-evidence variables X_i given its *Markov blanket*.
3. Wander randomly around the state space by flipping one variable at a time, but keeping the evidence fixed!

Markov Blanket

The Markov blanket of node (or variable) X consists of its parents, children, and children's parents (gray area)



$$P(X|mb(X)) = \alpha P(X|Parents(X)) \times \prod_{Y_i \in Children(X)} P(Y_i|Parents(Y_i))$$

Gibbs Sampling: Example

Query: $\mathbf{P}(\text{Rain} | \text{Splinker} = \text{true}, \text{WetGrass} = \text{true})$

1. state [Cloudy=true, Splinker=true, Rain=false, WetGrass=true]

Gibbs Sampling: Example

Query: $\mathbf{P}(\text{Rain} | \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$

1. state $[\text{Cloudy}=\text{true}, \text{Sprinkler}=\text{true}, \text{Rain}=\text{false}, \text{WetGrass}=\text{true}]$
2. *Cloudy* is sampled given its Markov blanket, i.e. we sample from $\mathbf{P}(\text{Cloudy} | \text{Sprinkler} = \text{true}, \text{Rain} = \text{false})$. If *Cloudy* = *false* then the new state is $[\text{false}, \text{true}, \text{false}, \text{true}]$

Gibbs Sampling: Example

Query: $\mathbf{P}(\text{Rain} | \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$

1. state $[\text{Cloudy}=\text{true}, \text{Sprinkler}=\text{true}, \text{Rain}=\text{false}, \text{WetGrass}=\text{true}]$
2. *Cloudy* is sampled given its Markov blanket, i.e. we sample from $\mathbf{P}(\text{Cloudy} | \text{Sprinkler} = \text{true}, \text{Rain} = \text{false})$. If *Cloudy* = *false* then the new state is $[\text{false}, \text{true}, \text{false}, \text{true}]$
3. *Rain* is sampled given its Markov blanket, i.e. we sample from $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{false}, \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$. If *Rain* = *true* then the new state is $[\text{false}, \text{true}, \text{true}, \text{true}]$

Gibbs Sampling: Example

Query: $\mathbf{P}(\text{Rain} | \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$

1. state $[\text{Cloudy}=\text{true}, \text{Sprinkler}=\text{true}, \text{Rain}=\text{false}, \text{WetGrass}=\text{true}]$
2. *Cloudy* is sampled given its Markov blanket, i.e. we sample from $\mathbf{P}(\text{Cloudy} | \text{Sprinkler} = \text{true}, \text{Rain} = \text{false})$. If *Cloudy* = *false* then the new state is $[\text{false}, \text{true}, \text{false}, \text{true}]$
3. *Rain* is sampled given its Markov blanket, i.e. we sample from $\mathbf{P}(\text{Rain} | \text{Cloudy} = \text{false}, \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$. If *Rain* = *true* then the new state is $[\text{false}, \text{true}, \text{true}, \text{true}]$
4. If this process observes 20 states with *Rain* = *true* and 60 states with *Rain* = *false*, then the answer to our example query is $\alpha < 20, 60 > = < 0.25, 0.75 >$

Gibbs Sampling: Algorithm

function GIBBS-ASK(X, \mathbf{e}, bn, N) **returns** an estimate of $\mathbf{P}(X|\mathbf{e})$
 local variables: \mathbf{N} , a vector of counts for each value of X , initially zero
 \mathbf{Z} , the nonevidence variables in bn
 \mathbf{x} , the current state of the network, initially copied from \mathbf{e}

 initialize \mathbf{x} with random values for the variables in \mathbf{Z}
 for $j = 1$ to N **do**
 for each Z_i in \mathbf{Z} **do**
 set the value of Z_i in \mathbf{x} by sampling from $\mathbf{P}(Z_i|mb(Z_i))$
 $\mathbf{N}[x] \leftarrow \mathbf{N}[x] + 1$ where x is the value of X in \mathbf{x}
 return NORMALIZE(\mathbf{N})

Keywords

- ▶ **Bayesian Network**
- ▶ **Approximate Inference**
- ▶ **Probabilistic Sampling**
- ▶ **Prior Sampling**
- ▶ **Rejection Sampling**
- ▶ **Estimated Probability Distribution**
- ▶ **Likelihood Weighting**
- ▶ **Markov Chain Monte Carlo**
- ▶ **Markov Blanket**
- ▶ **Gibbs Sampling**

Summary

- ▶ Probabilistic inference in Bayes nets means computing a probability distribution of a set of query variables, given a set of evidence and hidden variables
- ▶ Exact inference algorithms include inference by enumeration and variable elimination
- ▶ Approximate inference algorithms include rejection sampling, likelihood weighting, & Gibbs sampling
- ▶ Approximate inference can be used to answer queries in large Bayes—but the accuracy depends on the number of samples
- ▶ These algorithms combined with parameter learning (e.g. MLE) allow you to equip a system with probabilistic reasoning

Workshop:

- ▶ Exercises about probabilistic inference, Q&A
- ▶ Wednesday from 10:30-12hrs, room=INB1302

Reading of this week: Russell & Norvig, (2013). Artificial Intelligence: A Modern Approach, **chapter 14** (until Section 14.5)

Reading of next week: Russell & Norvig, (2013). Artificial Intelligence: A Modern Approach, **chapter 15** (with Nicola)